

Managing Class Imbalance

BA810: Supervised Machine Learning
Nachiketa Sahoo

Recap

1. Load and explore, training-test split
2. Consider data cleaning and preparation
 1. Imputation, feature selection, ...
3. Create pipelines
4. Build and evaluate **predictive models**
 1. Linear models, k-NN, SVM, DTree, ...
 2. Ensembles
5. Finetune hyperparameters of the most promising model
 1. Grid/Random search w. Halving
6. Estimate error on unused test-data

Ensembles

- Train multiple models (“weak learners”) and aggregate predictions (mean, mode, etc.)
 - Stabilizes prediction if models are uncorrelated
- Bagging (Bootstrap Aggregating)
 - Learn models from bootstrapped samples
 - Random forest: combine trees, each node considers random subset of columns
- Boosting: learn models sequentially
 - Prioritize records that were poorly predicted
- Voting: mode/mean predicted
 - Stacking: estimated scores/probability to outcome mapping is learnt
 - For some ranges of its score, the strongest model may be weaker than others (alone or combined)

Class Imbalance

- What is it?

- Uneven proportion of records from different classes
- Mild: 20–40%, Moderate: 1–20%, extreme: <1%

- What's the problem?

- Objective functions give each record equal weight

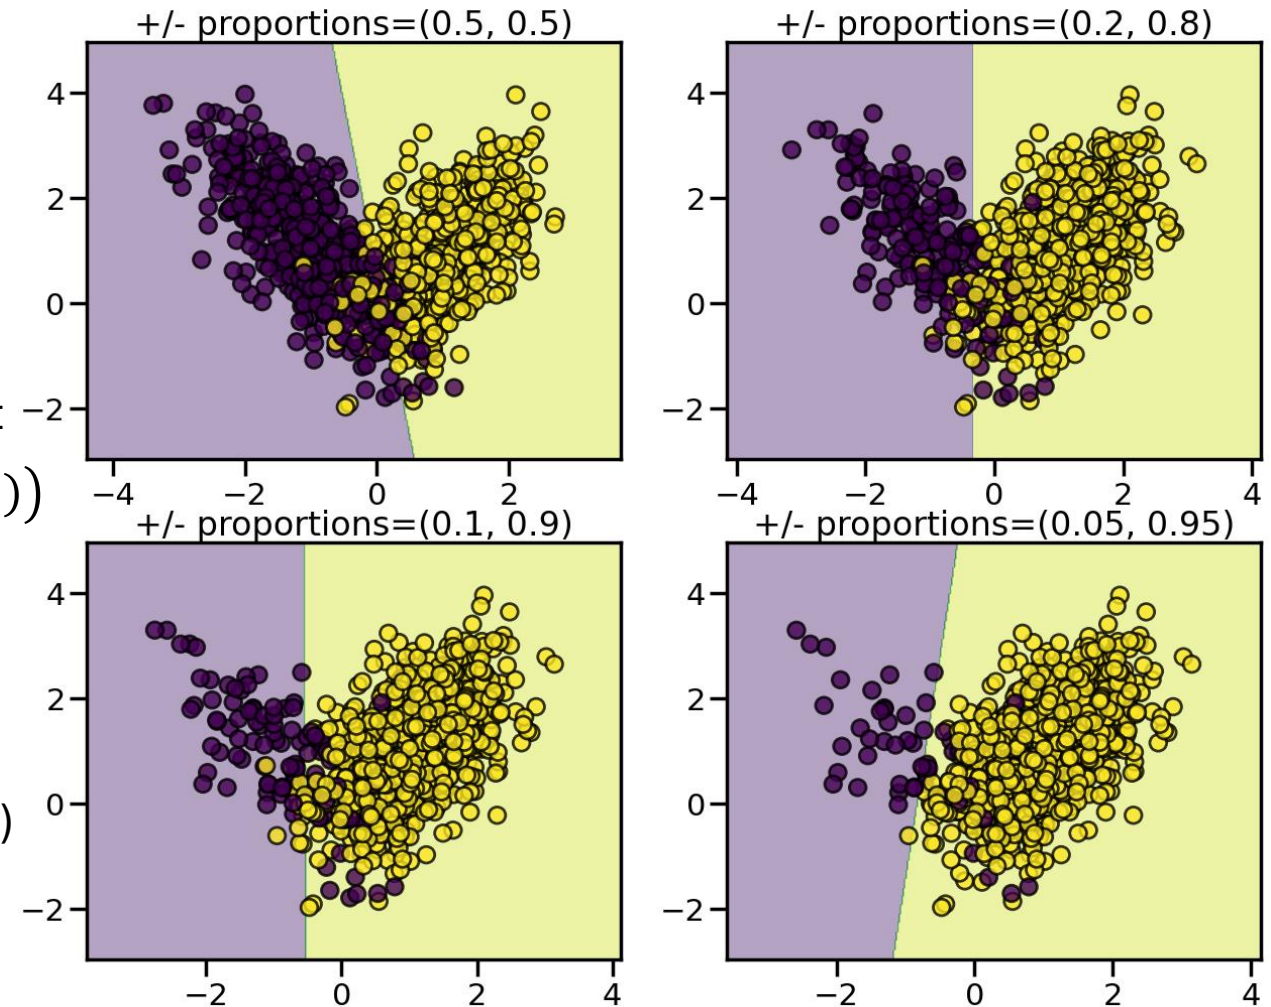
$$\min_{\beta} \sum_{i \in N_-} \text{Error}(y_i, f(x_i; \beta)) + \sum_{i \in N_+} \text{Error}(y_i, f(x_i; \beta))$$

where, N_- and N_+ are the sets of negative and positive training records, respectively

→ biases to predict majority class more often

- High accuracy but might predict poorly on the minority class (often of more interest, the positive)
 - Frauds, disease positive cases, etc.

Decision function of LogisticRegression



How to Handle?

Use the right metric

- Accuracy gives each *record* equal weight (thus much less weight to minority class)
 - $Accuracy = \frac{TN+TP}{TN+TP+FP+FN}$
 - Balanced accuracy gives each *class* equal weight
 - $\frac{1}{2} \left(\frac{TN}{TN+FP} + \frac{TP}{TP+FN} \right)$
- If you have cost of false negative vs false positive, use that!
 - See our lecture/labs on cost-sensitive evaluations
- These will highlight problems with default objective
 - Low balanced accuracy, high cost due to false negative,
...(though not fix them)

	PREDICTED CLASS		
		Positive	Negative
	Positive	TP	FN
ACTUAL CLASS	Negative	FP	TN

How to Handle?

Fix the objective

- If we care about records of minority class more, reflect that in the objective function


$$\sum_{i \in N_-} \text{Error}(y_i, f(x_i; \beta)) + w \sum_{i \in N_+} \text{Error}(y_i, f(x_i; \beta))$$

- Multiply by a class weight w : 5, 10, ...
 - Could be the cost of False Negative relative to False Positive
 - Or a value that gives both class equal weight in objective: $\left(\frac{|N_-|}{|N_+|}\right)$
- Easy to do; many models naturally support this—little added computational cost
- Generally, reduces accuracy, but improves balanced accuracy

How to Handle?

What if the model doesn't support w ? *Change Data*

- Oversample minority class: sample with replacement to make class sizes equal
 - Create duplicate minority records or copies with noise added (learn the model from modified data)

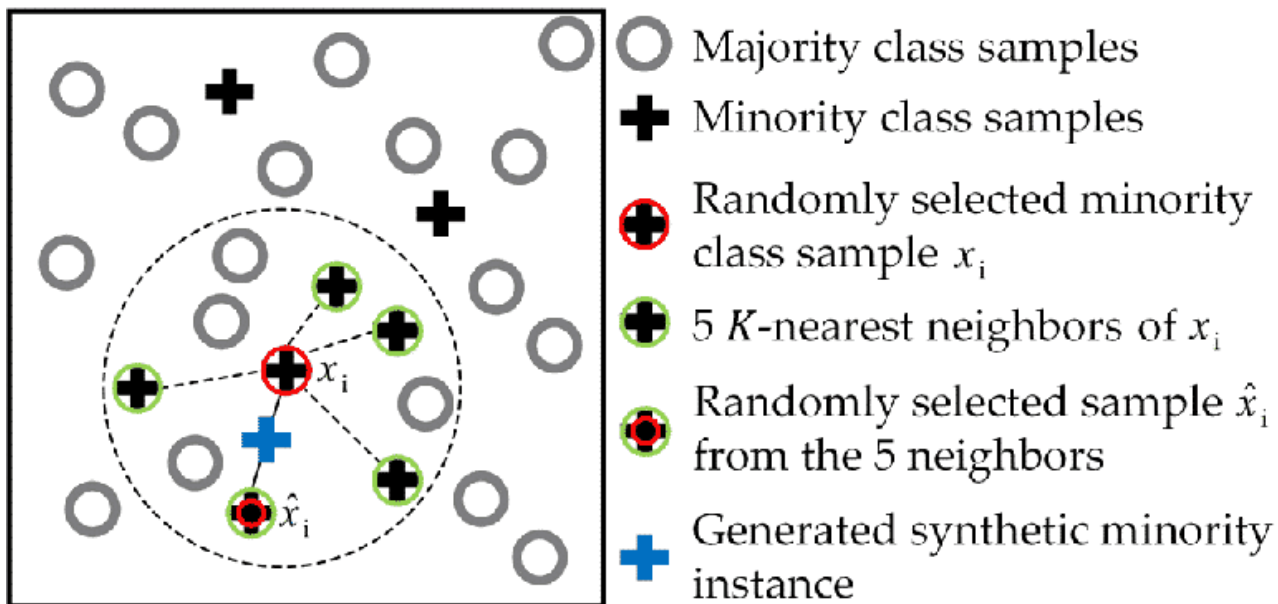


Based on example from <https://imbalanced-learn.org/>

How to Handle?

Change Data

- Synthetic Minority Over-sampling Technique: (SMOTE)




- Generate more around those points that are at the decision boundaries (ADASYN and SMOTE variants)


(source: https://rikunert.com/smote_explained)

Particularities of over-sampling with SMOTE and ADASYN


Decision function with SMOTE




Resampling with SMOTE



Decision function with ADASYN



Resampling with ADASYN



Based on example from <https://imbalanced-learn.org/>

How to Handle?

Other approaches

- When dataset is very large ... computing time and resource is the constraint
 - Under-sample the majority class
- Learn a standard classifier, but change the decision/probability threshold
 - Pick one that optimizes a desired objective (total cost, net benefit, or balanced accuracy, etc.)
 - Directly changes the location of decision boundary, countering the bias
 - See Lab 4
- Let's apply these in a default prediction problem with extreme class imbalance.
 - Same one used in Lab 4: 3.33% default and 96.67% don't.
 - Open [Lab 12](#).

Topics Recap

(With an eye on the final exam)

1. General Machine Learning concepts
 - Supervised/unsupervised, prediction vs causal inference, training/test split
2. Basic predictive models
 - Linear regression (R^2 , coefficient, p-value)
 - Classification (models, metrics, cost-based evaluation, ...)
 - Bias/variance, underfitting/overfitting
3. Model selection
 - Cross validation, regularization
 - Hyper-parameter search
4. End-to-end Machine Learning process
 - Preprocessing, use of pipelines
5. Specific predictive models
 - Support Vector Machines, Decision Tree, Ensembles
6. Managing imbalanced data in practice
 - Use of appropriate metrics, using class weight, over/under sample to balance

Tips for the Final Exam

- Prioritize lectures and labs, then textbook chapters
 - Don't worry about more complex math, but you should know the conceptual differences between methods and metrics (e.g., decision tree vs logistic regression, cross validation vs bootstrap, etc.)
 - Should know how to compute precision, recall, etc.
- Measurement metrics and interpretations are important
 - E.g., how to interpret logistic regression coefficient, or a confusion matrix
- Be precise in your answers to questions asking to explain (conceptual questions)
 - Stick to specified length limits